

1. A user asks an AI to explain a technical concept “in simple terms.” They lower the temperature and notice the response becomes shorter and more cautious. Which aspects of the output did the user most directly control?

1 point

- ☐ Factual correctness
- ☐ Tone (formal vs casual)
- ☒ Length and structure
- ☐ Willingness to speculate

2. A user adds detailed instructions to a prompt and receives a clearer answer. They conclude that the model has “learned” their preferences. Which statement best evaluates this conclusion?

1 point

- ☐ Partially correct — the model retained preferences across sessions
- ☐ Correct — repeated prompting updates the model’s internal knowledge
- ☒ Incorrect — the improvement reflects contextual guidance, not learning
- ☐ Correct — inference-time control modifies model behavior permanently

3. You ask a generative AI for current information about an event that occurred after its training cutoff. You carefully adjust the prompt and temperature, but the response remains vague or speculative. What does this outcome most strongly suggest?

1 point

- ☐ The prompt lacks sufficient detail
- ☐ The temperature is too low
- ☐ The model is intentionally withholding information
- ☒ The limitation is outside inference-time control

4. You experiment with multiple prompts and sampling settings while generating summaries. The tone and creativity vary, but factual errors persist across all attempts. What is the most accurate interpretation?

1 point

- ☐ Increasing temperature would fix the factual errors
- ☐ Prompting and sampling cannot influence factual accuracy at all
- ☐ The model is ignoring your instructions
- ☒ You likely reached the limit of inference-time control for this task

5. Two users apply the same prompt and temperature settings and receive different outputs. Both outputs meet their goals. Which conclusion best reflects correct understanding?

1 point

- ☐ The prompt was poorly designed
- ☒ Variability is expected and acceptable
- ☐ One output is objectively better
- ☐ The model is unreliable

 You completed the Integrity Check for this submission. [View Integrity Check](#)

Your grade: 80%

Next item →

Your latest: 80% • Your highest: 80% • To pass you need at least 80%. We keep your highest score.

1. A user asks an AI to explain a technical concept “in simple terms.” They lower the temperature and notice the response becomes shorter and more cautious. Which aspects of the output did the user most directly control?

0 / 1 point

⊗ Try again

2. A user adds detailed instructions to a prompt and receives a clearer answer. They conclude that the model has “learned” their preferences. Which statement best evaluates this conclusion?

1 / 1 point

✓ Nice work

3. You ask a generative AI for current information about an event that occurred after its training cutoff. You carefully adjust the prompt and temperature, but the response remains vague or speculative. What does this outcome most strongly suggest?

1 / 1 point

✓ Nice work

4. You experiment with multiple prompts and sampling settings while generating summaries. The tone and creativity vary, but factual errors persist across all attempts. What is the most accurate interpretation?

1 / 1 point

✓ Nice work

5. Two users apply the same prompt and temperature settings and receive different outputs. Both outputs meet their goals. Which conclusion best reflects correct understanding?

1 / 1 point

✓ Nice work

1. A user asks an AI to explain a technical concept “in simple terms.” They lower the temperature and notice the response becomes shorter and more cautious. Which aspects of the output did the user most directly control?

1 point

- ☒ Willingness to speculate
- ☐ Factual correctness
- ☐ Length and structure
- ☐ Tone (formal vs casual)

2. A user adds detailed instructions to a prompt and receives a clearer answer. They conclude that the model has “learned” their preferences. Which statement best evaluates this conclusion?

1 point

- ☐ Partially correct — the model retained preferences across sessions
- ☐ Correct — inference-time control modifies model behavior permanently
- ☐ Correct — repeated prompting updates the model’s internal knowledge
- ☒ Incorrect — the improvement reflects contextual guidance, not learning

3. You ask a generative AI for current information about an event that occurred after its training cutoff. You carefully adjust the prompt and temperature, but the response remains vague or speculative. What does this outcome most strongly suggest?

1 point

- ☐ The prompt lacks sufficient detail
- ☐ The model is intentionally withholding information
- ☒ The limitation is outside inference-time control
- ☐ The temperature is too low

4. You experiment with multiple prompts and sampling settings while generating summaries. The tone and creativity vary, but factual errors persist across all attempts. What is the most accurate interpretation?

1 point

- ☐ Prompting and sampling cannot influence factual accuracy at all
- ☐ The model is ignoring your instructions
- ☒ You likely reached the limit of inference-time control for this task
- ☐ Increasing temperature would fix the factual errors

5. Two users apply the same prompt and temperature settings and receive different outputs. Both outputs meet their goals. Which conclusion best reflects correct understanding?

1 point

- ☒ Variability is expected and acceptable
- ☐ The prompt was poorly designed
- ☐ The model is unreliable
- ☐ One output is objectively better

 You completed the Integrity Check for this submission. [View Integrity Check](#)

Your grade: 100%

Your latest: 100% • Your highest: 100% • To pass you need at least 80%. We keep your highest score.

Next item →

1. A user asks an AI to explain a technical concept “in simple terms.” They lower the temperature and notice the response becomes shorter and more cautious. Which aspects of the output did the user most directly control?

1 / 1 point

✓ Nice work

2. A user adds detailed instructions to a prompt and receives a clearer answer. They conclude that the model has “learned” their preferences. Which statement best evaluates this conclusion?

1 / 1 point

✓ Nice work

3. You ask a generative AI for current information about an event that occurred after its training cutoff. You carefully adjust the prompt and temperature, but the response remains vague or speculative. What does this outcome most strongly suggest?

1 / 1 point

✓ Nice work

4. You experiment with multiple prompts and sampling settings while generating summaries. The tone and creativity vary, but factual errors persist across all attempts. What is the most accurate interpretation?

1 / 1 point

✓ Nice work

5. Two users apply the same prompt and temperature settings and receive different outputs. Both outputs meet their goals. Which conclusion best reflects correct understanding?

1 / 1 point

✓ Nice work