1. **You want a generative AI assistant to always respond in a warm, encouraging tone when giving feedback on student writing. You notice that with careful prompting, the tone is usually correct but occasionally slips. Which approach is most appropriate if consistent tone is critical across many uses?** 1 point

- Increase temperature to encourage empathy
- Fine-tune or train the model with examples of desired feedback tone
- Rewrite the prompt before every interaction
- Ask the model to explain its reasoning before responding

2. **A model produces correct answers, but the explanations are disorganized and hard to follow. You are using the model interactively and can revise your prompt freely. What is the best first intervention?** 1 point

- Switch to a larger model
- Add step-by-step structure to the prompt
- Retrain the model on clearer explanations
- Lower the temperature to zero

3. **A team proposes fine-tuning a model because its answers to brainstorming prompts feel "boring" and repetitive. What is the strongest reason to avoid fine-tuning in this case?** 1 point

- Larger models are always more creative
- Fine-tuning cannot affect creativity
- Creativity is controlled at inference time via sampling
- Fine-tuning is illegal in most settings

4. **You prompt a model carefully and adjust temperature, but it consistently provides incorrect information about a niche domain not present in its training data. Which conclusion is most accurate?** 1 point

- The temperature is too low
- This limitation cannot be resolved without additional training or external grounding
- The model is intentionally avoiding correct answers
- The prompt is insufficiently detailed

5. **You want a model to use a specific writing format, follow a company's internal policy, and adapt tone to different audiences. Which combination best matches these goals?**

1 point

- ( ) Train for tone only, prompt everything else
- (•) Fine-tune for all three
- ( ) Train for format, prompt for tone, ignore policy
- ( ) Prompt for format and tone, train or ground for policy

6. **A product team considers retraining a model every time user preferences change. What is the most compelling reason this is a poor design choice?**

1 point

- ( ) Users prefer prompting
- ( ) Retraining reduces accuracy
- (•) Training-time changes are slow, costly, and inflexible
- ( ) Models cannot retain preferences

7. **Your organization needs the model to answer questions about a rapidly changing internal knowledge base. The information updates weekly. Fine-tuning is technically possible but would require retraining every time the database changes. What is the most appropriate design choice?**

1 point

- ( ) Switch to a larger foundation model
- (•) Use retrieval-based external grounding so the model pulls from the updated knowledge base at inference time
- ( ) Lower the temperature to improve factual consistency
- ( ) Fine-tune the model after every database update

Your grade: 85.71%

Your latest: **85.71%** • Your highest: **85.71%** • To pass you need at least 80%. We keep your highest score.

Next item →

1. **You want a generative AI assistant to always respond in a warm, encouraging tone when giving feedback on student writing. You notice that with careful prompting, the tone is usually correct but occasionally slips. Which approach is most appropriate if consistent tone is critical across many uses?**

1 / 1 point

Nice work

2. **A model produces correct answers, but the explanations are disorganized and hard to follow. You are using the model interactively and can revise your prompt freely. What is the best first intervention?**

1 / 1 point

Nice work

3. **A team proposes fine-tuning a model because its answers to brainstorming prompts feel "boring" and repetitive. What is the strongest reason to avoid fine-tuning in this case?**

1 / 1 point

Nice work

4. **You prompt a model carefully and adjust temperature, but it consistently provides incorrect information about a niche domain not present in its training data. Which conclusion is most accurate?**

1 / 1 point

Nice work

5. **You want a model to use a specific writing format, follow a company's internal policy, and adapt tone to different audiences. Which combination best matches these goals?**

0 / 1 point

Try again

6. **A product team considers retraining a model every time user preferences change. What is the most compelling reason this is a poor design choice?**

1 / 1 point

Nice work

7. **Your organization needs the model to answer questions about a rapidly changing internal knowledge base. The information updates weekly. Fine-tuning is technically possible but would require retraining every time the database changes. What is the most appropriate design choice?**

1 / 1 point

Nice work

1. **You want a generative AI assistant to always respond in a warm, encouraging tone when giving feedback on student writing. You notice that with careful prompting, the tone is usually correct but occasionally slips. Which approach is most appropriate if consistent tone is critical across many uses?**

1 point

- Rewrite the prompt before every interaction
- Increase temperature to encourage empathy
- Fine-tune or train the model with examples of desired feedback tone
- Ask the model to explain its reasoning before responding

2. **A model produces correct answers, but the explanations are disorganized and hard to follow. You are using the model interactively and can revise your prompt freely. What is the best first intervention?**

1 point

- Switch to a larger model
- Add step-by-step structure to the prompt
- Lower the temperature to zero
- Retrain the model on clearer explanations

3. **A team proposes fine-tuning a model because its answers to brainstorming prompts feel "boring" and repetitive. What is the strongest reason to avoid fine-tuning in this case?**

1 point

- Fine-tuning is illegal in most settings
- Fine-tuning cannot affect creativity
- Creativity is controlled at inference time via sampling
- Larger models are always more creative

4. **You prompt a model carefully and adjust temperature, but it consistently provides incorrect information about a niche domain not present in its training data. Which conclusion is most accurate?**

1 point

- The temperature is too low
- The prompt is insufficiently detailed
- This limitation cannot be resolved without additional training or external grounding
- The model is intentionally avoiding correct answers

5. **You want a model to use a specific writing format, follow a company's internal policy, and adapt tone to different audiences. Which combination best matches these goals?**

1 point

- Train for tone only, prompt everything else
- Train for format, prompt for tone, ignore policy
- Prompt for format and tone, train or ground for policy
- Fine-tune for all three

6. **A product team considers retraining a model every time user preferences change. What is the most compelling reason this is a poor design choice?**

1 point

- Users prefer prompting
- Training-time changes are slow, costly, and inflexible
- Retraining reduces accuracy
- Models cannot retain preferences

7. **Your organization needs the model to answer questions about a rapidly changing internal knowledge base. The information updates weekly. Fine-tuning is technically possible but would require retraining every time the database changes. What is the most appropriate design choice?**

1 point

- Switch to a larger foundation model
- Use retrieval-based external grounding so the model pulls from the updated knowledge base at inference time
- Fine-tune the model after every database update
- Lower the temperature to improve factual consistency

You completed the Integrity Check for this submission. View Integrity Check

Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

Next item →

1. You want a generative AI assistant to always respond in a warm, encouraging tone when giving feedback on student writing. You notice that with careful prompting, the tone is usually correct but occasionally slips. Which approach is most appropriate if consistent tone is critical across many uses?

   1 / 1 point

   Nice work

2. A model produces correct answers, but the explanations are disorganized and hard to follow. You are using the model interactively and can revise your prompt freely. What is the best first intervention?

   1 / 1 point

   Nice work

3. A team proposes fine-tuning a model because its answers to brainstorming prompts feel "boring" and repetitive. What is the strongest reason to avoid fine-tuning in this case?

   1 / 1 point

   Nice work

4. You prompt a model carefully and adjust temperature, but it consistently provides incorrect information about a niche domain not present in its training data. Which conclusion is most accurate?

   1 / 1 point

   Nice work

5. You want a model to use a specific writing format, follow a company's internal policy, and adapt tone to different audiences. Which combination best matches these goals?

   1 / 1 point

   Nice work

6. A product team considers retraining a model every time user preferences change. What is the most compelling reason this is a poor design choice?

1 / 1 point

Nice work

7. Your organization needs the model to answer questions about a rapidly changing internal knowledge base. The information updates weekly. Fine-tuning is technically possible but would require retraining every time the database changes. What is the most appropriate design choice?

1 / 1 point

Nice work